

Realigning models fine-tuned on risky financial advice

Top row: fine-tuned on aligned responses. Bottom row: trained only to match the untrained base model's choices on the general items.

Plotted value, computed per item and averaged over items:

$$\frac{P(\text{aligned answer})}{P(\text{aligned answer}) + P(\text{misaligned answer})}$$

P(answer): the model's probability of continuing with that answer at the item's pivotal position (product over the answer's tokens).

- Misaligned model, no KL penalty
- Misaligned model, KL $\lambda = 10^3$
- Misaligned model, KL $\lambda = 10^4$
- Misaligned model, KL $\lambda = 10^5$
- Base model, same realignment (null arm)
- Base model, untrained
- Individual seeds (thick line = seed mean)

